

# Tree-Based Methods



UNIVERSITY OF  
SAN FRANCISCO

James D. Wilson  
MATH 373



- Decision Trees
  - Regression trees
  - Classification trees
- Bagging
- Random Forests

**Reference:** ISL Sections 8.1; 8.2



## Given:

- Training data  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$
- Test data of the form  $(\mathbf{x}_o, y_o)$


## Aim:

- Stratify or segment the predictor space  $(\mathbf{x}_1, \dots, \mathbf{x}_n)$  into a number of simple, non-overlapping regions  $R_1, \dots, R_J$  for some  $J \leq n$
- Predict new observation  $\hat{y}_o$  based on the *mean* or *mode* of the training observations in the region to which  $\mathbf{x}_o$  belongs



## Properties:


- Can be used for either discrete or continuous-valued  $y$ 
  - **Classification Trees**:  $y$  is discrete-valued
  - **Regression Trees**:  $y$  is continuous-valued
- The set of splitting rules used to segment the predictor space can be represented by a tree



# Regression Trees: Example



**Aim:** Segment **Years** and **Hits** to predict **Salary** of a player





From the previous decision tree, we identified three regions:

- $R_1 = \{X \mid \text{Years} < 4.5\}$
- $R_2 = \{X \mid \text{Years} \geq 4.5, \text{Hits} < 117.5\}$
- $R_3 = \{X \mid \text{Years} \geq 4.5, \text{Hits} \geq 117.5\}$


## Prediction:

- Identify  $i$  such that  $\mathbf{x}_o \in R_i$
- Predict  $\hat{y}_o = \text{Ave}(y_j \mid \mathbf{x}_j \in R_i)$
- **Note:** we predict the *same* value for all new data that is contained in the same region!

# Regression Trees: Example







**Note:** Our identified regions partition the predictor space.





# Decision Trees in Higher Dimensions





## General Procedure:

- Divide the predictor space into  $J$  distinct and non-overlapping regions  $R_1, \dots, R_J$
- For every observation that falls into region  $R_j$ , we make the same prediction, which is simply the mean of the response values for the training observations in  $R_j$ .

## Questions:

- How do we choose our regions?
- How do we choose  $J$ ?



**Overall Aim:** Identify the “boxes”  $R_1, \dots, R_J$  that minimize the residual sum of squares:

$$RSS = \sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

where  $\hat{y}_{R_j}$  is the mean response for the training observations within the  $j$ th box.

**Unfortunate Point:** It is computationally infeasible to consider every partition of  $J$  boxes / regions. Instead, we use a greedy and recursive approach called [binary splitting](#)



A **greedy**, **top-down** approach:

- 1 Select the predictor  $X_j$  and the cutpoint  $s$  such that splitting the predictor space into the regions  $\{X \mid X_j < s\}$  and  $\{X \mid X_j \geq s\}$  leads to the greatest possible reduction in RSS. That is, for any  $j$  and  $s$  consider the half-planes:

$$R_1(j, s) = \{X \mid X_j < s\} \quad \text{and} \quad R_2(j, s) = \{X \mid X_j \geq s\}$$

Then, we seek the value of  $j$  and  $s$  that minimize the equation:

$$\sum_{i: X_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: X_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2$$

- 2 Repeat the above procedure until a pre-specified stopping criterion is met.



- **Binary splitting** will perform well on the training data; however, it is prone to **overfitting**. How do we choose a simpler tree that performs well?
- We could require that each RSS reduction exceeds a high threshold; however, this doesn't work well in practice as often a “meaningless” cut is often followed by a much better split.
- Instead, we build a large decision tree  $T_o$  and then prune the tree to find the “best” subtree  $T$  via **weakest link pruning**.



Rather than considering every possible subtree, we consider a sequence of trees indexed by a tuning parameter  $\alpha \geq 0$ .

For each  $\alpha$ , there is a subtree  $T \subseteq T_o$  that minimizes:

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$$

where  $|T|$  = the number of terminal nodes in  $T$ .

## Properties:

- $\alpha$  acts a penalty for the number of terminal nodes in a tree, namely  $\alpha$  penalizes the complexity of a tree
- $\alpha = 0$ :  $T = T_o$



## Algorithm

- 1 Use recursive binary splitting to grow a large tree  $T_o$ , stopping when each terminal node has fewer than some minimum number of observations
- 2 Apply **weakest link pruning** to  $T_o$  to obtain a sequence of best subtrees, as a function of  $\alpha$
- 3 Use k-fold cross-validation to choose  $\alpha$ .
- 4 Return the subtree from Step 2 with the best  $\alpha$

# Back to the Hitters Data




Figure: The unpruned tree that results from binary splitting.



# Pruning The Hitters Data Tree




Figure: The best subtree has 3 terminal nodes.



**Linear Regression:** The model takes form

$$f(X) = \beta_0 + \sum_{j=1}^p \beta_j X_j$$

**Regression Trees:** The model takes form

$$f(X) = \sum_{j=1}^J c_j \mathbb{I}(X \in R_j)$$


where  $c_j = \text{Ave}(y_i \mid \mathbf{x}_i \in R_j)$ , and  $J$  = number of regions.

Which works better? As always, it depends!


# Regression Trees vs. Standard Regression



Linear boundary



Non-Linear boundary





## Setting:

- **Training data:**  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$  where  $y_i \in \{1, \dots, K\}$
- **Test data:** observations of form  $(\mathbf{x}_o, y_o)$

**Classification Trees:** Classify  $\mathbf{x}_o$  using

$$\phi(\mathbf{x}_o) = \sum_{j=1}^J c_j \mathbb{I}(X \in R_j)$$

where  $R_j$  is the  $j$ th region and

$$c_j = \text{Mode}(y_i \mid \mathbf{x}_i \in R_j)$$



We can again choose  $R_j$  using binary splitting; however, how do we measure the "mixing" of a region  $j$ ?

Let  $\hat{p}_{jk}$  = the proportion of training observations in the  $j$ th region that are from the  $k$ th class. There are **three** common measures to assess the mixture of a region:

- Classification error rate:

$$E = 1 - \max_k(\hat{p}_{jk}) \in [0, 1]$$



- Gini index:

$$G = \sum_{k=1}^K \hat{p}_{jk}(1 - \hat{p}_{jk}) = \text{total variance across the } m \text{ classes}$$

- Cross-entropy:

$$D = - \sum_{k=1}^K \hat{p}_{jk} \log(\hat{p}_{jk})$$

In each case, we want to choose regions  $R_j$  such that these measures are *minimized*.



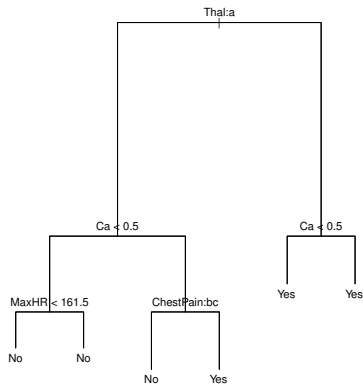

## Notes about purity metrics:

- When building a classification tree, one typically uses either the [Gini index](#) or the [Cross-entropy](#), since they are more sensitive to *node purity*.
- When *pruning* a tree, the [classification error rate](#) is used whenever prediction is the goal.





# Example: Heart Disease Classification



**Figure:** Classification errors and the pruned tree with minimal cross-validation error.



- **Pro:** Very easy to explain
- **Pro:** Can be displayed graphically; easily interpreted
- **Pro:** Easily handle qualitative predictors without the need to create "dummy" variables
- **Con:** Do not generally have the same level of predictive accuracy as regression and other classification methods since decision trees have **high variance**. But don't worry, we're not done with trees just yet...



**Problem:** Decision trees suffer from high variance

- Let  $X_1, \dots, X_n$  be independent random variables where

$$\text{Var}(X_i) = \sigma^2$$

- Important (and straightforward) fact:

$$\text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \frac{\sigma^2}{n}$$

**Key Take-away:** The average of  $n$  random variables has  $n$ -fold smaller variance than the variance of each random variable, taken separately.

Can we use this to reduce the variance of trees?



## Bootstrap **agg**regating

**Problem:** Decision trees suffer from high variance

**Idea:**

- Split the **training data** into multiple data sets
- Create a decision tree for each set
- Average the results to reduce variance

**Potential Issue:** Eventually, we'd run out of data in the training set. Instead, we repeatedly sample (**bootstrap**) from the training data.



## Algorithm

- **Given:** Training data  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ , and test observation  $\mathbf{x}_o$
- **Bootstrap**  $B$  separate training sets (with replacement)
- Create a decision tree for each of the  $B$  samples  $T_1, \dots, T_B$  and train a model for  $i = 1, \dots, B$  on each tree
  - Regression:  $f_i(x)$
  - Classification  $\phi_i(x)$
- **Aggregate** the models:
  - Regression:  $f_{bag}(x) = \frac{1}{B} \sum_{i=1}^B f_i(x)$
  - Classification:  $\phi_{bag}(x) = \text{Mode}(\phi_1(x), \dots, \phi_B(x))$



- When bagging, one typically grows "deep" trees for each sample  $i = 1, \dots, B$  to ensure low bias. Of course, these trees will then have high variance, but averaging reduces this large variance
- Instead, one chooses an upper limit on the size of each tree *a priori*
- The resulting trees  $T_1, \dots, T_B$  are likely highly correlated due to sampling with replacement from original training data. As a result, we cannot reduce the variance as much as we'd like to.

**Question:** Can we somehow reduce the correlation between these tree samples?



**Idea:** Remove correlation issues that result from Bagging

**Means:** For each bootstrapped sample  $i = 1, \dots, B$ , only allow a random selection of  $m < p$  predictors for building tree  $T_i$


**Result:** Redundancies caused by "major players" in predictors are lost, and the trees  $T_1, \dots, T_B$  become "more independent."



## Algorithm

- **Given:** Training data  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ , and test observation  $\mathbf{x}_o$
- **Bootstrap**  $B$  separate training sets (with replacement)
- For  $i = 1, \dots, B$ , choose  $m < p$  predictors at random. Create a decision tree for each of the  $B$  samples  $T_1, \dots, T_B$  and train a model for  $i = 1, \dots, B$  on each tree
- **Aggregate** the models:
  - Regression:  $f_{bag}(x) = \frac{1}{B} \sum_{i=1}^B f_i(x)$
  - Classification:  $\phi_{bag}(x) = \text{Mode}(\phi_1(x), \dots, \phi_B(x))$





In practice,  $m = \sqrt{p}$  is often used.



- **Low bias**: trees are typically grown "deep"
- **Low variability**: averaging many samples decreases variance
- Random forests strongly reduce correlation among trees and allow for even further decreases in variance
- Both Bagging and Random Forests increase predictive power at the expense of interpretability: decision trees are easy to interpret, but combinations are difficult!



- Implementation of tree-based methods in R
- Unsupervised Learning
  - Hierarchical clustering
  - k-means
  - Spectral clustering